

CSc 360: Operating Systems (Spring 2015)

Assignment 2: A Priority Queueing System (PQS)

Due Date: Feb. 27, 2015

1 Introduction

In this assignment, you need to face the second programming challenge: implementing a task scheduler. You will learn how to use the three programming constructs provided by the POSIX `pthread` library:

1. threads
2. mutexes
3. condition variables (convars).

Your goal is to simulate a priority queueing system (PQS). In such a system, there are arriving customers and a clerk who serves the customers. When a customer arrives and the clerk is busy with serving other customers, the arriving customer needs to wait. We assume a work-conserving system, i.e., the clerk cannot be idle if there are customers waiting for service. Each customer is given a priority number. To simplify your job, we assume customers are given before hand and their description is stored in a file (refer to Section 2.2).

You need to use threads to simulate the customers arriving and waiting for service, and your program will schedule the services to the customers to meet the requirements defined in Section 2.3.

2 Customers

2.1 Properties of Customers

Each customer, which will be simulated by a thread, has the following attributes:

1. **Arrival Time:** It indicates when the customer will arrive.
2. **Service Time:** It indicates the time required for the clerk to serve the customer (i.e., from the time when the clerk starts serving the customer to the time when the service is over). We assume **non-preemptive** scheduling, meaning that when a customer is being served, other customers with a higher priority cannot interrupt the service.
3. **Priority:** It is an integer number between 1 (inclusive) and 10 (inclusive), indicating the priority of the customer. **A larger number means a higher priority.**

All times are measured in 10ths of a second. The times will be simulated by having your threads, which represent customers, to call `usleep()` for the required amount of time.

2.2 The Format of Input File

Your program (named PQS) will accept *one* parameter on the command line:

`PQS customers.txt`

where `customers.txt` is the name of the input file.

2.2.1 File Format

The input file is a text file and has a simple format. The first line contains the number of customers that will be simulated. After that, each line contains the information about a single customer, such that:

1. The first character specifies the unique number of the customer.
2. A colon(:) immediately follows the unique number of the customer.
3. Immediately following is an integer that indicates the **arrival time** of the customer.
4. A comma(,) immediately follows the previous number.
5. Immediately following is an integer that indicates the **service time** of the customer.
6. A comma(,) immediately follows the previous number.
7. Immediately following is an integer that indicates the **priority** of the customer.
8. A newline (\n) ends a line.

To not kill yourself by checking the false input file, you *should* build a correct input file for your own test. We will not test your code with an input file that does not comply with the above format.

2.2.2 An Example

The following file specifies three customers.

```
3
1:3,60,3
2:6,70,1
3:3,50,3
```

The customer specifications are listed in the following table:

| customer No. | Arrival time | Service time | Priority |
|--------------|--------------|--------------|----------|
| 1 | 0.3s | 6s | 3 |
| 2 | 0.6s | 7s | 1 |
| 3 | 0.3s | 5s | 3 |

2.3 Simulation Rules

The rules enforced by the PQS are:

1. Only one customer is on service at any given time.
2. When a customer arrives, the clerk serves the arriving customer if no other customer is being served; otherwise, the arriving customer must wait in the queue.
3. When multiple customers arrive at the same time and no other customer is being served, the arriving customer with the highest priority will be served first. If they have the same priority, the one that has the shortest service time will be served first. If there is still a tie, the one that appears first in the input file will be served first.
4. When the clerk finishes the service of a customer, all waiting customers compete for service according to the following rules:
 - (a) The one with the highest priority will be served first.
 - (b) If there is a tie at the highest priority, the one whose arrival time is the earliest will be served first.
 - (c) If there is still a tie, the one who has the shortest service time will be served first.
 - (d) If there is still a tie, the one who appears first in the input file will be served first.

2.4 Output

Your simulation is required to output all events and state changes showing the internal behavior of PQS. The messages must include, but are not limited to:

1. A customer arrives
2. A customer waits for the completion of another customer who is being served.
3. The clerk starts serving a customer
4. The clerk finishes the service to a customer

You must:

1. print the arrival of each customer using the following format string to show the customer attributes:

```
"customer %2d arrives: arrival time (%.2f), service time (%.1f), priority (%2d). \n"
```

2. print the ID of the waiting customer and the ID of the customer being served with the format string:

```
"customer %2d waits for the finish of customer %2d. \n"
```

3. print the time when the clerk starts serving a customer using the format string:

```
"The clerk starts serving customer %2d at time %.2f. \n"
```

4. print the time when the clerk finishes its service to a customer using the format string:

```
"The clerk finishes the service to customer %2d at time %d. \n"
```

Note that the output of times (including arrival time, the time when the clerk starts serving a customer, the time when the clerk finishes the service of a customer) is **relative machine time**, calculated by the machine time when the output event occurs minus the machine time when the simulation starts. Therefore, the output of the times may not exactly matches (but should be close to) the results with manual calculation.

3 Submission: Deliverable 1 (Design Due: 23:59 pm, Feb 17, 2013)

You will write a design document which answers the following questions. It is recommended that you think through the following questions *very carefully* before answering them.

Unlike Assignment 1, debugging will be harder after the basic design has been coded. Therefore, it is very important to ensure that the basic design is correct. So think about the following carefully and then write down the answers.

1. How many threads are you going to use? Specify the task that you intend each thread to perform.
2. Do the threads work independently? Or, is there an overall "controller" thread?
3. How many mutexes are you going to use? Specify the operation that each mutex will guard.
4. Will the main thread be idle? If not, what will it be doing?
5. How are you going to represent customers? what type of data structure will you use?
6. How are you going to ensure that data structures in your program will not be modified concurrently?
7. How many convars are you going to use? For each convar:
 - (a) Describe the condition that the convar will represent.
 - (b) Which mutex is associated with the convar? Why?

(c) What operation should be performed once `pthread_cond_wait()` has been unblocked *and* re-acquired the mutex?

8. In 25 lines or less, briefly sketch the overall algorithm you will use. You may use sentences such as:

If the service to a customer is finished, release trans mutex.

Note: You will submit the design with `connex`, and you are required to type in your document and submit it in pdf file format. Other file format or handwriting will **not** be accepted.

4 Submission: Deliverable 2 (Code Due: 11:59 pm, Feb. 27, 2015)

The code is submitted through `connex`. The tutorial instructor will give the detailed instruction in the tutorial.

4.1 Submission Requirements

1. The name of the submission file must be `p2.tar.gz`
2. `p2.tar.gz` must contain all your files in a directory named `p2`
3. Inside the directory `p2`, there must be a `Makefile`.
4. Invoking `make` on it must result in an executable named `PQS` being built, *without user intervention*.
5. You may *not* submit the assignment with a compiled executable and/or object (`.o`) files.
6. Inside the directory `p2`, there must be an input file following the format described in Section 2.2, although we will test you code using our own input file.

5 Plagiarism

This assignment is to be done individually. You are encouraged to discuss the design of the solution with your classmates, but each student must implement their own assignment. The markers may submit your code to an automated plagiarism detection service.

6 Miscellaneous

6.1 Manual Pages

Be sure to study the `man` pages for the various functions to be used in the assignment. For example, the `man` page for `pthread_create` can be found by typing the command:

```
$ man pthread_create
```

At the end of this assignment you should be at least familiar with the following functions:

1. File access functions:

- (a) `atoi`
- (b) `fopen`
- (c) `feof`
- (d) `fgetc` and `fgets`
- (e) `fclose`

2. Thread creation functions:

- 137 (a) `pthread_create`
- 138 (b) `pthread_exit`
- 139 (c) `pthread_join`

140 3. Mutex manipulation functions:

- 141 (a) `pthread_mutex_init`
- 142 (b) `pthread_mutex_lock`
- 143 (c) `pthread_mutex_unlock`

144 4. Condition variable manipulation functions:

- 145 (a) `pthread_cond_init`
- 146 (b) `pthread_cond_wait`
- 147 (c) `pthread_cond_broadcast`
- 148 (d) `pthread_cond_signal`

149 It is absolutely critical that you read the `man` pages, and attend the tutorials.
150 Your best source of information, as always, is the `man` pages.
151 For help with POSIX threads:

152 <http://www.opengroup.org/onlinepubs/007908799/xsh/pthread.h.html>

153 A good overview of `pthread` can be found at: <http://www.llnl.gov/computing/tutorials/pthreads/>

154 6.2 Important Notes

155 We want to (re-)emphasize the following points:

- 156 1. **You are required to type in your design document. Hand writing will not be accepted.**
- 157 2. We will give a **time quota of 2 minutes** for your program to run on a given input. This time quota is given
158 so that non-terminating programs can be killed. So make sure your input file does not include too many long
159 customers (e.g., arrive too late or the service time too long). Since your program simulates customers in 10ths
160 of a second, this should not be an issue, at all.
- 161 3. **It is required that you use relative machine time. This is to avoid cheating with an implemen-**
162 **tation that does not really simulate the customers but instead performs an offline analysis to**
163 **obtain scheduling results. The markers will read your C code to ensure that the `pthread` library**
164 **is used as required. Offline analysis means that your program does not simulate mutual exclusion**
165 **and thread synchronization but obtains the output based on algorithm analysis. You will get 0**
166 **marks if you are caught using offline analysis.**
- 167 4. As you progress through your degree the projects and assignments will continue to become more complicated
168 and difficult. It is impossible to describe in detail every possible case/feature in the assignment specification.
169 Instead you are expected to apply the techniques you have learned so far in your degree, use common sense, and
170 ask questions to lecture instructor or TA when something is unclear. We will announce further clarification, if
171 necessary, on [Connex](#). Complaining the specification is unclear at the last minute is **not acceptable**.
- 172 5. You are required to use C. Any other language is **not acceptable**.
- 173 6. You are required to strictly follow the input file format. Failing to do so will result in the deduction of scores.
- 174 7. You should use the Linux machines in ECS242 to test your work.
- 175 8. Programming with semaphore is permitted but not recommended. You are warned that debugging with
176 semaphore is much harder.

7 Marking

We will mark your design document and your code submission.

7.1 Design Document (10% of this assignment)

You are required to answer all questions listed in Section 3. Your design will be marked based on the clear and correct logic and the correctness of your algorithm.

7.2 Code Submission(90% of this assignment)

7.2.1 Functionality

1. Your **PQS** must correctly schedule the customer services, with our own test files. We will not disclose all test files before the final submission.
2. You are required to catch return errors of important function calls, especially when a return error may result in the logic error or malfunctioning of your program.
3. Your program must output at least the information described in Section 2.4.

7.2.2 Code Quality

We cannot specify completely the coding style that we would like to see but it includes the following:

1. Proper decomposition of a program into subroutines (and multiple source code files when necessary)—A 1000 line C program as a single routine would fail this criterion.
2. Comment—judiciously, but not profusely. Comments also serve to help a marker, in addition to yourself. To further elaborate:
 - (a) Your favorite quote from Star Wars or Douglas Adams' Hitch-hiker's Guide to the Galaxy does not count as comments. In fact, they simply count as anti-comments, and will result in a loss of marks.
 - (b) Comment your code in English. It is the official language of this university.
3. Proper variable names—**leia** is not a good variable name, it never was and never will be.
4. Small number of global variables, if any. Most programs need a very small number of global variables, if any. (If you have a global variable named **temp**, think again.)
5. **The return values from all system calls and function calls listed in the assignment specification should be checked and all values should be dealt with appropriately.**

If you are in doubt about how to write good C code, you can easily find

<http://www.chris-lott.org/resources/cstyle/>.

The <http://www.chris-lott.org/resources/cstyle/indhill-annot.html> is an excellent short style guide.

7.3 Detailed Test Plan

The detailed test plan for the code submission is as follows.

| Components | Weight |
|-------------------------------------|--------|
| Make file | 1 |
| Input file | 1 |
| Normal cases | 4 |
| Priority tie | 2 |
| Arrival time tie | 2 |
| Transmission time tie | 2 |
| All tie | 2 |
| Special cases (illegal values) | 1 |
| Output format | 1 |
| Catch system call return values | 1 |
| Comments (functional decomposition) | 2 |
| Code style | 1 |
| Critical sections | 2 |
| Readme | 1 |
| Total Weight | 23 |

209 Note that a score of 23/23 means that you obtain 90% (Section 7.2) of this assignment. If that is the case,
 210 congratulations!

211 Last but not the least, please read this document **carefully**. If you are not sure, confirm with the course
 212 instructor.